



MAXIMIZE INSIGHTS AND PERFORMANCE WITH THE RIGHT TIME-SERIES VIEW


RAJNISH KUMAR
DATA ARCHRTECT



Content

CAGG

Materialized View in PostgreSQL



1 Understand different types of views

2 Views Vs Materialized view

3 Materialized View Pros and Cons

4 CAGG – Continuous Aggregate view

5 View Vs Materialized View Vs CAGG

6 Conclusion

Different Types Of View



VIEW



Live, Real-Time
Queries



MATERIALIZED VIEW



Static Snapshot
on Refresh



**CONTINUOUS
AGGREGATE
(CAGG)**



Auto-Updated
Incrementally



Best for
Time-Series Data

What is a View?

A **view** is a **virtual table** in a database.

It **does not store data physically**; instead, it stores a **saved SQL query**.

How a View Works

when you query a view, the database:

1. Takes the SQL query stored in the view
2. Executes it on the underlying base tables
3. Returns the result **as if it were a real table**

For User's perspective – it behaves just like a table.

What is a Materialized View?

A **materialized view** is a **physical snapshot** of a query's result.

Unlike a standard view (which stores only a query), a materialized view **stores the actual data on disk**.

How Materialized View Works

1. Creation

- The database executes the query **once**
- The result set is stored as a **physical table behind the scenes**

How To Create,

```
CREATE MATERIALIZED VIEW view_user_daily_activity AS
SELECT date_trunc('month', order_date) AS month,
       SUM(steps) AS total_step
FROM device_data_table
GROUP BY 1;
```

2. Access

- Querying a materialized view **reads precomputed data**
- No joins or aggregations are recalculated at query time

```
SELECT * FROM view_user_daily_activity;
```

This is why it's **fast**.

3. Refresh

- Materialized views **do not update automatically**
- They remain static until refreshed

Postgre:

```
REFRESH MATERIALIZED VIEW CONCURRENTLY view_user_daily_activity;
```

NOTE: avoids blocking reads, we need to create required index

What **CONCURRENTLY** does

It allows **reads to continue while the materialized view is being refreshed**.

Without it:

- PostgreSQL **locks the materialized view**
- All **SELECT queries wait or fail**
- Bad for dashboards & live apps

With it:

- Users keep querying the **old data**
- PostgreSQL builds the **new data in the background**
- Once done, it **atomically swaps old → new**

Quick Comparison: View vs Materialized View

Feature	View	Materialized View
Data Stored	✗ No	✓ Yes
Storage	Minimal (query only)	Uses disk space
Freshness	Always fresh	Stale until refreshed
Query Speed	Depends on query complexity	Very fast ⚡
Use Case	Abstraction, security, simplicity	Analytics, reporting, heavy queries

View = live query shortcut * Materialized View = precomputed snapshot

Pros & Cons of Materialized Views



Pros

- ✓ Speed
- ✓ Offloading
- ✓ Predictable Query Time
- ✓ Fast services



Cons

- ✗ Data Staleness
- ✗ Disk Space Usage
- ✗ Manual Refresh Needed
- ✗ Expensive Refresh



Continuous Aggregate (CAGG)

A **Continuous Aggregate (CAGG)** is a **materialized view** in TimescaleDB that is **continuously and incrementally updated** as new data is added to a **hypertable**. Unlike a regular materialized view that requires manual refreshes, a continuous aggregate automatically keeps aggregated data up-to-date with minimal overhead.


Key points:

- Works on **hypertables** (TimescaleDB's optimized time-series tables).
- Performs **incremental refreshes**, meaning only new or changed data is aggregated.
- Can **downsample data**, e.g., daily, weekly, monthly from a high-frequency time-series table.
- Improves **query performance** for large datasets.

Continuous Aggregates (CAGGS) in TimescaleDB

Powerful Time-Series Data Aggregation

- ✓ Incremental updates as new data arrives.
- ✓ Improved query performance for analytical workloads.
- ✓ Automatic refresh policies (e.g., daily, monthly).
- ✓ Downsample & summarize high-frequency data easily



```
CREATE MATERIALIZED VIEW cagg_user_daily_activity
WITH (timescaledb.continuous) AS
SELECT
    time_bucket('1 month', order_date) AS month,
    SUM(steps) AS total_step
FROM device_data_table
GROUP BY month;
```

WITH (**timescaledb.continuous**)

- Marks this as a **Continuous Aggregate** that refreshes automatically and incrementally.

Refresh Policy

You can also set a **refresh interval**, so TimescaleDB updates the aggregate automatically:

```
SELECT add_continuous_aggregate_policy('cagg_user_daily_activity',  
    start_offset => INTERVAL '1 month',  
    end_offset => INTERVAL '1 day',  
    schedule_interval => INTERVAL '1 hour');
```

- `start_offset` → how far back to start refreshing.
- `end_offset` → leave small buffer for late-arriving data.
- `schedule_interval` → how often TimescaleDB updates the CAGG.

Continuous Aggregate (CAGG)

AUTO & INCREMENTAL UPDATE



Regular Materialized View (MV)

MANUAL & FULL REFRESH



Conclusion:

Continuous Aggregates (CAGGs) in TimescaleDB provide a **powerful and efficient way to aggregate time-series data**. Unlike regular materialized views that require manual refreshes and can be slow on large datasets, CAGGs **incrementally update** as new data arrives, significantly improving query performance.

By using features like **time_bucket** and **automatic refresh policies**, developers can **downsample, summarize, and analyze high-frequency data** (daily, monthly, or yearly) without worrying about stale results or heavy computation.



THANK YOU
FOR YOUR TIME AND ATTENTION

RAJNISH KUMAR,
Data Architect

